Statistical Inference with R

Linear & Logistic Regression

•••

GW Libraries Workshop
Dan Kerchner ~ March 4, 2022

go.gwu.edu/rstats

Super-Brief Review of Inference for Regression

High-Level Objective

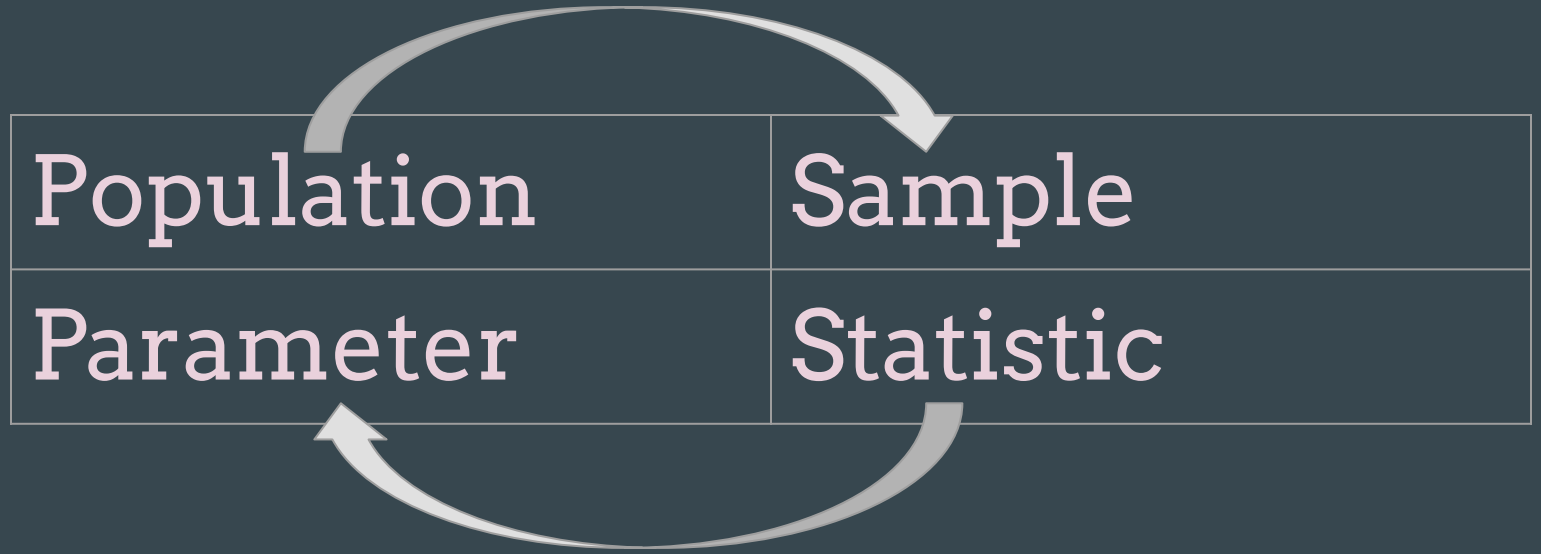

Regression Model

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \ldots + \beta_n X_n$$

(and an interaction term might look like $\beta_{12} X_1 X_2$)

Interpretation

β_0 = Mean Y when X_i values are 0

β_i = Mean change in Y for a 1-point increase in X_i , adjusting for other X variables

Correlation between dependent & independent variables

Pearson correlation (ϱ) measures strength and direction (+/-) of linear association between X_i and Y. Ranges from -1 to 1.

If relationship looks non-linear (but monotonic) then **Spearman** correlation should be considered.

$H_0: \varrho = 0$
$H_1: \varrho \neq 0$

Valid if joint distribution of X, Y is bivariate normal

Linear Regression Model

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \ldots + \beta_n X_n$$

where Y is a continuous outcome

Interpretation

β_0 = Mean Y when X_i values are 0

β_i = Mean change in Y for a 1-point increase in X_i, adjusting for other X variables

Linear Regression Assumptions

- Observations are independent
- Linearity
- Homoscedasticity
- Normality

GLMs - Generalized Linear Models

$$g(\mu) = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \ldots + \beta_n X_n$$

where $\mu = E(Y)$

Common link functions

$$g(\mu) = \mu$$

$$g(\mu) = \log(\mu)$$

$$g(\mu) = \log(\mu / (1 - \mu)) = \text{logit}(\mu)$$

GLMs: (Binary) Logistic Regression Model

odds

$$log(\frac{p}{1-p}) = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \ldots + \beta_n X_n$$

where p = Probability of Y = 1

1-p = Probability of Y = 0

Interpretation

β_i = log OR (Odds Ratio) for having $Y = 1$ for a 1-point increase in X_i, *adjusting for other predictors*

e^{β_i} = OR for having $Y = 1$ for a 1-point increase in X_i

Binary Logistic Regression Assumptions

- Binary outcome (0, 1)
- Each predictor is linearly related to the log odds of the outcome

Inference for Regression Modeling

Confidence Interval

95% CI for β = (0.44, 0.49)

Hypothesis Testing

H_0: $\beta = \beta_0$ ← Null Hypothesis
H_A: $\beta \neq \beta_0$ ← Alternative Hypothesis

p-value: Chance that we are rejecting H_0 when we should not be

Goals

Today's Goal

- Learn to use R to read in data and conduct regression analysis and associated inference tests
  - Checking assumptions
  - Visualizing
  - Computing p-values, regression coefficients, confidence intervals, and odds ratios (for logistic models)

Today: 2 Scenarios

- Linear Regression (continuous outcome)
  - Single variable (1 continuous predictor)
  - Multivariable (1 continuous, 1 categorical predictor)
- Logistic Regression (categorical outcome)
  - Multivariable (continuous and categorical predictors)

Today's Data Set: Siddarth et al., 2018

Siddarth P, Burggren AC, Eyre HA, Small GW, Merrill DA (2018) Sedentary behavior associated with reduced medial temporal lobe thickness in middle-aged and older adults. PLoS ONE 13(4): e0195549. doi.org/10.1371/journal.pone.0195549

- Examined associations between sedentary behavior and medial temporal lobe (MTL) subregion integrity
- 35 non-demented middle-aged and older adults
- Measured physical activity levels w/questionnaire
- Measured MTL thickness w/MRI scan
- Adjusted for age

Today's Data Set: Framingham Heart Study

- framinghamheartstudy.org
- Long-term prospective study of the etiology of cardiovascular disease among a population of subjects in Framingham, MA
- Began in 1948 with 5,209 subjects
- Is the source of the term "risk factor"
- Over 3,000 peer-reviewed papers published based on this study
- Participants were each followed for a total of 24 years for cardiovascular events (heart attack, stroke, death, etc.)

Some Handy R Links

Some R packages & functions for regression++

- **`lme4`** - Linear mixed-effects models
- **`MASS`** - Model selection
- **`ts(), arima()`** - Time series object, model (part of **`stats`**)
- **`forecast`** - Forecasting for time series and linear models
- **`mice`** - imputation of missing data
- **`medflex`** - mediation analysis
- **`splines2`** - regression spline basis functions
- **`survival`** - survival analysis; JM - joint modeling
- AND MANY, MANY, MANY MORE

Tutorials

- RStudio R paths: education.rstudio.com/learn/
- Data Carpentry & Software Carpentry:
  - datacarpentry.org and software-carpentry.org
- Linkedin Learning @ GW: go.gwu.edu/linkedinlearning
- r-tutor.com/r-introduction & r-tutor.com/elementary-statistics
- UCLA Data Analysis Examples: stats.idre.ucla.edu/other/dae/
- Visualizing regression results: worldbank.github.io/r-econ-visual-library/RegressionCoef.html
- R Graph Gallery (w/code): r-graph-gallery.com

Books you can access for free

- Free books online - Hadley Wickham:
  - R for Data Science r4ds.had.co.nz
  - Advanced R adv-r.hadley.nz/
- Through your GW library privileges:

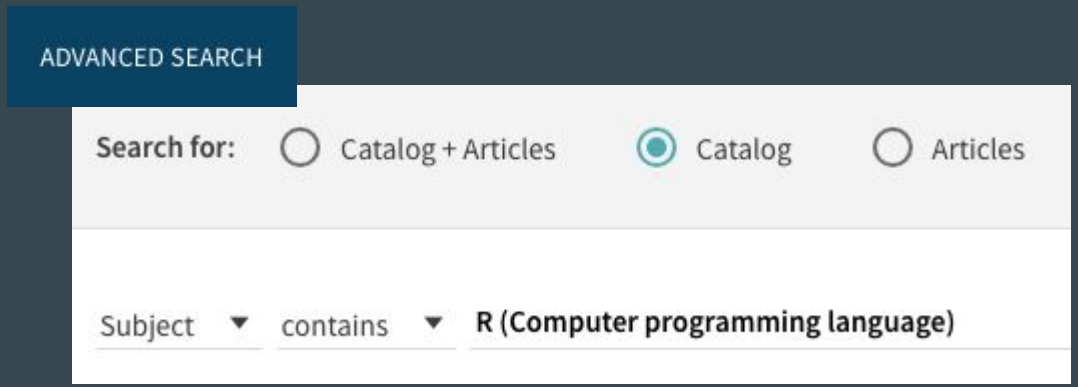

Reference Links

- R language (CRAN): r-project.org
- R search engine: rseek.org
- rstudio.com
  - Cheat Sheets! rstudio.com/resources/cheatsheets
- stackoverflow.com

Statistics+R help @ GW

R-Statistics Appointments:
academiccommons.gwu.edu/data-consulting

Also...

Appointments with me:
calendly.com/kerchner

Coding consultations (Python, R, git, etc.):
calendly.com/gwul-coding/

Thanks!

Dan Kerchner kerchner@gwu.edu